

Statistical Thinking in Python Part 2

5). Putting it all together: a case study

a). EDA of beak depths of Darwin's finches

Create bee swarm plot

```
_ = sns.swarmplot(x='year', y='beak_depth', data=df)
```


Label the axes

```
_ = plt.xlabel('year')
```

```
_ = plt.ylabel('beak depth (mm)')
```

Show the plot

```
plt.show()
```



b). ECDFs of beak depths

```
# Compute ECDFs
```

```
x_1975, y_1975 = ecdf(bd_1975)
```

```
x_2012, y_2012 = ecdf(bd_2012)
```

```
# Plot the ECDFs
```

```
_ = plt.plot(x_1975, y_1975, marker='.', linestyle='none')
```

```
_ = plt.plot(x_2012, y_2012, marker='.', linestyle='none')
```

```
# Set margins
```

```
plt.margins(0.02)
```

```
# Add axis labels and legend
```


```
_ = plt.xlabel('beak depth (mm)')
```

```
_ = plt.ylabel('ECDF')
```

```
_ = plt.legend(('1975', '2012'), loc='lower right')
```

```
# Show the plot
```

```
plt.show()
```



c). Parameter estimates of beak depths

```
# Compute the difference of the sample means: mean_diff
mean_diff = np.mean(bd_2012) - np.mean(bd_1975)

# Get bootstrap replicates of means
bs_replicates_1975 = draw_bs_reps(bd_1975, np.mean, size=10000)
bs_replicates_2012 = draw_bs_reps(bd_2012, np.mean, size=10000)

# Compute samples of difference of means: bs_diff_replicates
bs_diff_replicates = bs_replicates_2012 - bs_replicates_1975

# Compute 95% confidence interval: conf_int
conf_int = np.percentile(bs_diff_replicates, [2.5, 97.5])

# Print the results
print('difference of means =', mean_diff, 'mm')
print('95% confidence interval =', conf_int, 'mm')
```

<script.py> output:

difference of means = 0.226220472441 mm

95% confidence interval = [0.05633521 0.39190544] mm

d). Hypothesis test: Are beaks deeper in 2012?

```
# Compute mean of combined data set: combined_mean
combined_mean = np.mean(np.concatenate((bd_1975, bd_2012)))

# Shift the samples
bd_1975_shifted = (bd_1975 - np.mean(bd_1975))+combined_mean
bd_2012_shifted = (bd_2012 - np.mean(bd_2012))+combined_mean

# Get bootstrap replicates of shifted data sets
bs_replicates_1975 = draw_bs_reps(bd_1975_shifted, np.mean, 10000)
bs_replicates_2012 = draw_bs_reps(bd_2012_shifted, np.mean, 10000)

# Compute replicates of difference of means: bs_diff_replicates
bs_diff_replicates = bs_replicates_2012 - bs_replicates_1975

# Compute the p-value
p = np.sum(bs_diff_replicates >= mean_diff) / len(bs_diff_replicates)

# Print p-value
print('p =', p)
```

<script.py> output:

p = 0.0034

e). EDA of beak length and depth**# Make scatter plot of 1975 data**

```
_ = plt.plot(bl_1975, bd_1975, marker='.',  
            linestyle='none', color='blue', alpha=0.5)
```

Make scatter plot of 2012 data


```
_ = plt.plot(bl_2012, bd_2012, marker='.',  
            linestyle='none', color='red', alpha=0.5)
```

Label axes and make legend

```
_ = plt.xlabel('beak length (mm)')  
_ = plt.ylabel('beak depth (mm)')  
_ = plt.legend(('1975', '2012'), loc='upper left')
```

Show the plot

```
plt.show()
```



f). Linear Regression

```
# Compute the linear regressions
```

```
slope_1975, intercept_1975 = np.polyfit(bl_1975, bd_1975, 1)
```

```
slope_2012, intercept_2012 = np.polyfit(bl_2012, bd_2012, 1)
```

```
# Perform pairs bootstrap for the linear regressions
```

```
bs_slope_reps_1975, bs_intercept_reps_1975 = \
    draw_bs_pairs_linreg(bl_1975, bd_1975, 1000)
```

```
bs_slope_reps_2012, bs_intercept_reps_2012 = \
    draw_bs_pairs_linreg(bl_2012, bd_2012, 1000)
```

```
# Compute confidence intervals of slopes
```

```
slope_conf_int_1975 = np.percentile(bs_slope_reps_1975,[2.5,97.5])
```

```
slope_conf_int_2012 = np.percentile(bs_slope_reps_2012,[2.5,97.5])
```

```
intercept_conf_int_1975 = np.percentile(bs_intercept_reps_1975,[2.5,97.5])
```

```
intercept_conf_int_2012 = np.percentile(bs_intercept_reps_2012,[2.5,97.5])
```

```
# Print the results
```

```
print('1975: slope =', slope_1975,
      'conf int =', slope_conf_int_1975)
print('1975: intercept =', intercept_1975,
      'conf int =', intercept_conf_int_1975)
print('2012: slope =', slope_2012,
      'conf int =', slope_conf_int_2012)
print('2012: intercept =', intercept_2012,
      'conf int =', intercept_conf_int_2012)
```

```
<script.py> output:
```

```
1975: slope = 0.465205169161 conf int = [ 0.33851226  0.59306491]
```

```
1975: intercept = 2.39087523658 conf int = [ 0.64892945  4.18037063]
```

```
2012: slope = 0.462630358835 conf int = [ 0.33137479  0.60695527]
```

```
2012: intercept = 2.97724749824 conf int = [ 1.06792753  4.70599387]
```

g). Displaying the linear regression result**# Make scatter plot of 1975 data**

```
_ = plt.plot(bl_1975, bd_1975, marker='.',
            linestyle='none', color='blue', alpha=0.5)
```

Make scatter plot of 2012 data

```
_ = plt.plot(bl_2012, bd_2012, marker='.',
            linestyle='none', color='red', alpha=0.5)
```

Label axes and make legend

```
_ = plt.xlabel('beak length (mm)')
_ = plt.ylabel('beak depth (mm)')
_ = plt.legend(('1975', '2012'), loc='upper left')
```

Generate x-values for bootstrap lines: x


```
x = np.array([10, 17])
```

Plot the bootstrap lines

```
for i in range(100):
    plt.plot(x, bs_slope_reps_1975[i]*x+bs_intercept_reps_1975[i],
            linewidth=0.5, alpha=0.2, color='blue')
    plt.plot(x, bs_slope_reps_2012[i]*x+bs_intercept_reps_2012[i],
            linewidth=0.5, alpha=0.2, color='red')
```

Draw the plot again

```
plt.show()
```



h). Beak length to depth ratio

```
# Compute length-to-depth ratios
```

```
ratio_1975 = bl_1975/ bd_1975
```

```
ratio_2012 = bl_2012/ bd_2012
```

```
# Compute means
```

```
mean_ratio_1975 = np.mean(ratio_1975)
```

```
mean_ratio_2012 = np.mean(ratio_2012)
```

```
# Generate bootstrap replicates of the means
```

```
bs_replicates_1975 = draw_bs_reps(ratio_1975, np.mean, size=10000)
```

```
bs_replicates_2012 = draw_bs_reps(ratio_2012, np.mean, size=10000)
```

```
# Compute the 99% confidence intervals
```

```
conf_int_1975 = np.percentile(bs_replicates_1975, [0.5,99.5])
```

```
conf_int_2012 = np.percentile(bs_replicates_2012, [0.5,99.5])
```

```
# Print the results
```

```
print('1975: mean ratio =', mean_ratio_1975,
```

```
      'conf int =', conf_int_1975)
```

```
print('2012: mean ratio =', mean_ratio_2012,
```

```
      'conf int =', conf_int_2012)
```

<script.py> output:

```
1975: mean ratio = 1.57888237719 conf int = [ 1.55668803  1.60073509]
```

```
2012: mean ratio = 1.46583422768 conf int = [ 1.44363932  1.48729149]
```


i). EDA od hieratibility**# Make scatter plots**

```
_ = plt.plot(bd_parent_fortis, bd_offspring_fortis,  
            marker='.', linestyle='none', color='blue', alpha=0.5)  
_ = plt.plot(bd_parent_scandens, bd_offspring_scandens,  
            marker='.', linestyle='none', color='red', alpha=0.5)
```

Label axes

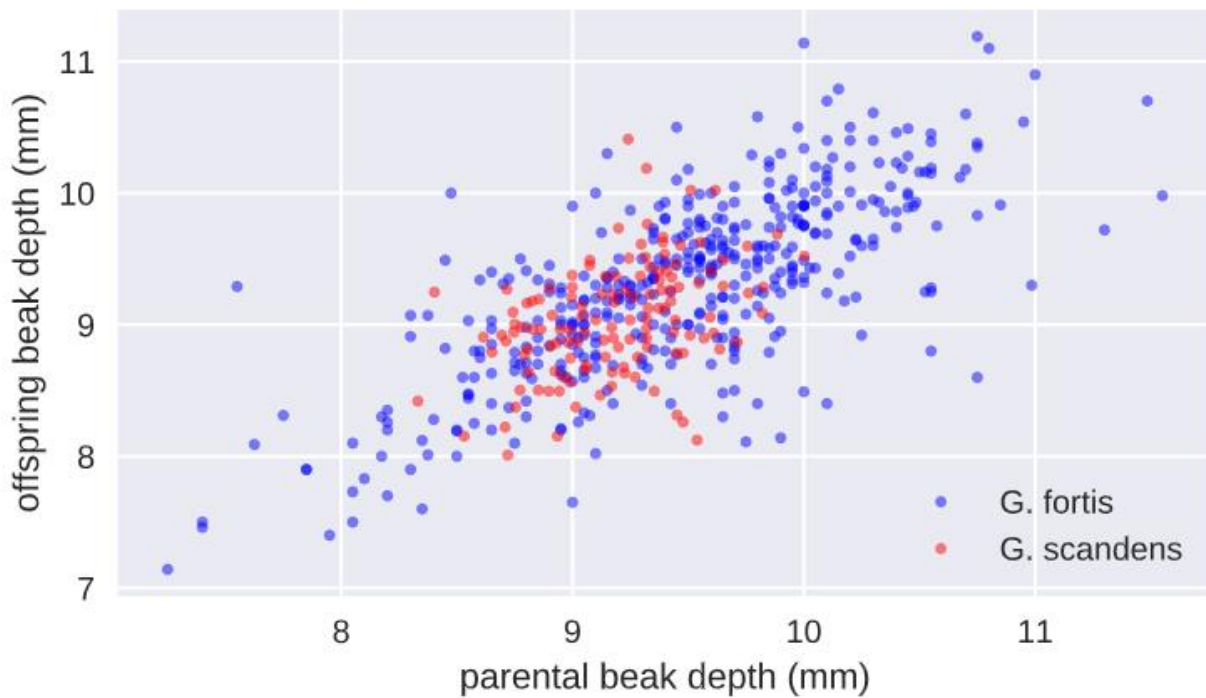
```
_ = plt.xlabel('parental beak depth (mm)')  
_ = plt.ylabel('offspring beak depth (mm)')
```

Add legend

```
_ = plt.legend(('G. fortis', 'G. scandens'), loc='lower right')
```

Show plot

```
plt.show()
```



j). Correlation of offspring and Parental data

```
def draw_bs_pairs(x, y, func, size=1):  
    """Perform pairs bootstrap for single statistic."""  
  
    # Set up array of indices to sample from: inds  
    inds = np.arange(len(x))  
  
    # Initialize replicates: bs_replicates  
    bs_replicates = np.empty(size)  
  
    # Generate replicates  
    for i in range(size):  
        bs_inds = np.random.choice(inds, len(inds))  
        bs_x, bs_y = x[bs_inds], y[bs_inds]  
        bs_replicates[i] = func(bs_x, bs_y)  
  
    return bs_replicates
```

k). Pearson correlation of offspring and parental data

```
# Compute the Pearson correlation coefficients
```

```
r_scandens = pearson_r(bd_parent_scandens, bd_offspring_scandens)
```

```
r_fortis = pearson_r(bd_parent_fortis, bd_offspring_fortis)
```

```
# Acquire 1000 bootstrap replicates of Pearson r
```

```
bs_replicates_scandens = draw_bs_pairs(bd_parent_scandens, bd_offspring_scandens, pearson_r, 1000)
```

```
bs_replicates_fortis = draw_bs_pairs(bd_parent_fortis, bd_offspring_fortis, pearson_r, 1000)
```

```
# Compute 95% confidence intervals
```

```
conf_int_scandens = np.percentile(bs_replicates_scandens,[2.5,97.5])
```

```
conf_int_fortis = np.percentile(bs_replicates_fortis,[2.5,97.5])
```

```
# Print results
```

```
print('G. scandens:', r_scandens, conf_int_scandens)
```

```
print('G. fortis:', r_fortis, conf_int_fortis)
```

<script.py> output:

G. scandens: 0.41170636294 [0.26564228 0.54388972]

G. fortis: 0.728341239552 [0.6694112 0.77840616]

1). Measuring Heritability

```
def heritability(parents, offspring):
```

```
    """Compute the heritability from parent and offspring samples."""
```

```
    covariance_matrix = np.cov(parents, offspring)
```

```
    return covariance_matrix[0,1] / covariance_matrix[0,0]
```

```
# Compute the heritability
```

```
heritability_scandens = heritability(bd_parent_scandens, bd_offspring_scandens)
```

```
heritability_fortis = heritability(bd_parent_fortis, bd_offspring_fortis)
```

```
# Acquire 1000 bootstrap replicates of heritability
```

```
replicates_scandens = draw_bs_pairs(bd_parent_scandens, bd_offspring_scandens, heritability, 1000)
```

```
replicates_fortis = draw_bs_pairs(bd_parent_fortis, bd_offspring_fortis, heritability, 1000)
```

```
# Compute 95% confidence intervals
```

```
conf_int_scandens = np.percentile(replicates_scandens, [2.5,97.5])
```

```
conf_int_fortis = np.percentile(replicates_fortis, [2.5,97.5])
```

```
# Print results
```

```
print('G. scandens:', heritability_scandens, conf_int_scandens)
```

```
print('G. fortis:', heritability_fortis, conf_int_fortis)
```

<script.py> output:

```
G. scandens: 0.548534086869 [ 0.34395487  0.75638267]
```

```
G. fortis: 0.722905191144 [ 0.64655013  0.79688342]
```

j). Is beak depth heritable at all in *G. scandens*?

```
# Initialize array of replicates: perm_replicates
perm_replicates = np.empty(10000)

# Draw replicates
for i in range(10000):
    # Permute parent beak depths
    bd_parent_permuted = np.random.permutation(bd_parent_scandens)
    perm_replicates[i] = heritability(bd_parent_permuted, bd_offspring_scandens)

# Compute p-value: p
p = np.sum(perm_replicates >= heritability_scandens) / len(perm_replicates)

# Print the p-value
print('p-val =', p)
```

<script.py> output:

p-val = 0.0